


java.interview.preparation

This blog is to provide details for java interview preparation. Everybody is busy in day to day project work and most of times we are not able to spend time due to workload. I am trying to provide the consolidated details of java topics and sharing my java interview experiences.

Home Concurrency Design Pattern Spring MVC Spring JMS Data structure Collection Threading Hibernate Face 2 Fa
Interview Planning Interview Questions Memory Mgmt & GC XML Puzzles DB Java 1.5 Transaction Inner Class Cache

Database Interview Questions for Java Developers Part 2

Here are more DB interview question for java developer. Mostly I have been asked question preliminary on DB Index, Query optimisation, Standard query and DB functions. DB Interview Questions : [Part -1](#).



Q 1- How to Enable/Disable Indexes?

A - Disable Index

ALTER INDEX [IndexName] ON TableName DISABLE GO

Enable Index

ALTER INDEX [IndexName] ON TableName [REBUILD](#)

Q 2- What is max number of columns in index?

A- 32

Q 3- How many maximum columns in partitions can be used?

A - 16

Q 4- How to replace specific values from query result? i.e. how query return 'M' and 'F' can be converted into 'Male' and 'Female'.

A - This can be done using Oracle [Decode](#) function the syntax will be like

```
SELECT employee_name,
DECODE(type, 'M', 'Male',
        'F', 'Female') type
FROM employee;
```

Q 5 - If the column of the table is having case sensitive data like 'Abc' and 'ABC' and we want to do case sensitive search on table. How we can improve performance of case sensitive search?

A- In this case we can define case sensitive index on table so that each search is restricted to particular value compare to full table scan. You can force all your data to be the same case by using UPPER() or LOWER():

```
select * from my_table where upper(column_1) = upper('my_string');
```

or

```
select * from my_table where lower(column_1) = lower('my_string');
```

If column_1 is not indexed on upper(column_1) or lower(column_1), as appropriate, this may force a full table scan. In order to avoid this you can create a function-based index.

```
create index my_index on my_table ( lower(column_1) );
```

If you're using LIKE then you have to concatenate a % around the string you're searching for.

```
select * from my_table where lower(column_1) = lower('my_string') || '%';
```

Q 5 - What are Synonyms?

A Synonyms give you the ability to provide alternate names for [database objects](#). You can [alias](#) object names; for example, using the Employee table as Emp. You can also shorten names. This is especially useful when dealing with three and four part names; for example, shortening server.database.owner.object to object.

Q 6 - What is the Difference between Update Lock and Exclusive Lock?

A -When Exclusive Lock is on any process, no other lock can be placed on that row or table. Every other process have to wait till Exclusive Lock completes its tasks.

Update Lock is a type of Exclusive Lock, except that it can be placed on the row which already has Shared Lock on it. Update Lock reads the data of the row which has the Shared Lock as soon as the Update Lock is ready to change the data it converts itself to the Exclusive Lock.

Q 7 -What is a Surrogate Key?

A - A surrogate key is a substitution for the natural primary key. It is just a unique identifier or number for each row that can be used for the primary key to the table. The only requirement for a surrogate primary key is that it should [be unique](#) for each row in the table. It is useful because the natural

INFOLINKS

BEST OFF



Best Java

[Java, Jav](#)
Raleigh, N
Computer
[ServiceNc](#)
Southamp
Sandersor
[Scientific](#)
München,
Informate
[Senior .NE](#)
Broadview
MCG - Mi
[Senior Jav](#)
Arnhem, C
Amoria Bc
[Java Devi](#)
Dublin, Du
Connectu
[Hadoop J](#)
Darmstad
Cognizant
[C++ Deve](#)
Chicago, I
TeamBrac
[Senior Jav](#)
Waalwijk,
Ardekay
[Mobile Ap](#)
Phoenix, A
USG
[POST A J](#)

MUST REA

primary key can change and this makes updates more difficult. Surrogated keys are always integer or numeric.

A 8 - Design the Database process which can handle 1 million requests at same time and What should be DB consideration?

A - All the major databases and operating systems will work in 1 million requests at same time; provided sufficient hardware with sufficient bandwidth to memory, disk, and network are provided.

All databases are built for this kind of scenario -- that is, where you need to update lots of rows from different clients all the time. This is what indexes and primary keys are used for, and the database is optimized for updating in this manner. (i.e., UPDATE your_table SET your_column=your_value where your_key=12)

The hardware requirements are going to be your biggest issue, and I suspect you will need to think about quite a lot here, including:

- Failover (what happens when your main server fails?)
- Clustering (You may simply need to have more than one database server to handle your load)
- Processors (2? 4? 8? 16? Multi-core? Does the db provide multi-core support that is well optimized?)
- Memory (The faster the better, but your chipset must be able to handle the bandwidth as well)
- Disk (Faster I/O the better. eSATA/SATA or Fiber, etc.)
- Network (You'll need lots of bandwidth to handle this kind of traffic)
- Virtualization (Does it make sense to build this as real hardware? Why not as virtual servers in the cloud? Amazon / Rackspace / etc.?)

Thankfully a good majority of the scaling issues are handled either at the hardware or db level. That is, you can start your database now on a slow server and as traffic increases you can adjust accordingly with minimal changes to your code or your db structure. What will change is your hardware and your db configuration. Most database servers will support failover/clustering with little to no change to your existing code or structures. (But be sure to research possible costs and the efficiency thereof first)

Good discussion on same issue.

Q 9 - Remove duplicates from table?

A:

```
DELETE FROM TableName
WHERE ID NOT IN (SELECT MAX(ID)
                 FROM   TableName
                 GROUP BY Column1,
                        Column2,
                        Column3
                 HAVING MAX(ID) IS NOT NULL)
```

Q 10 : Finding N'th Maximum salary SQL query?

```
// Using Sub query
FROM Employee Emp1
WHERE (N-1) = (
    SELECT COUNT(DISTINCT(Emp2.Salary))
    FROM Employee Emp2
    WHERE Emp2.Salary > Emp1.Salary
)
```

SELECT *

```
//Another way to get 2'nd maximum salary
Select max(Salary) From Employee e where
e.sal < ( select max(sal) from employee );
```

Q 11 - Database Query Optimisation tips with JDBC:

A : Use prepared statement pooling

Database supports prepared statement pooling for pooled connections, as discussed in the JDBC 3.0 specification, through the TimesTenObservableConnectionDS class. Note that statement pooling is transparent to an application. Use of the PreparedStatement object, including preparing and closing the statement, is no different.

Use arrays of parameters for batch execution

You can improve performance by using groups, referred to as *batches*, of statement executions, calling the addBatch() and executeBatch() methods for Statement or PreparedStatement objects.


A batch can consist of a set of INSERT, UPDATE, DELETE, or MERGE statements. Statements that return result sets, such as SELECT statements, are not allowed in a batch. A SQL statement is added to a batch by calling addBatch() on the statement object. The set of SQL statements associated with a batch are executed through the executeBatch() method.

For PreparedStatement objects, a batch consists of repeated executions of a statement using different input parameter values. For each set of input values, create the batch by using appropriate setXXX() calls followed by the addBatch() call. The batch is executed by the executeBatch() method.

```
// turn off autocommit
conn.setAutoCommit(false);
Statement stmt = conn.createStatement();
stmt.addBatch("INSERT INTO employees VALUES (1000, 'Joe Jones')");
stmt.addBatch("INSERT INTO departments VALUES (260, 'Shoe')");
stmt.addBatch("INSERT INTO emp_dept VALUES (1000, 260)");
// submit a batch of update commands for execution
int[] updateCounts = stmt.executeBatch();
conn.commit();
```

Bulk fetch rows of data

Oracle provides an extension that allows an application to fetch multiple rows of data. For applications that retrieve large amounts of data, fetching multiple rows can increase performance greatly. However, when using Read Committed isolation level, locks are held on all rows being retrieved until the application has received all the data, decreasing concurrency. For more information on this feature, see "Fetching multiple rows of data".



POPULAR I

Unix comma
interview qu

Java Memor

Hibernate In

Java Interv

Data Structu
developer

Database In
Developers

Java 1.7 fea

JMS Interv

Design Patter

java interview

FACEBOOK



Like P




GOOGLE+

FOLLOWE

Join 1
with Google

Members (



Already a me

TOTAL PA



ABOUT ME

Java Int

Use the ResultSet method getString() sparingly

Because Java strings are immutable, the ResultSet method getString() must allocate space for a new string in addition to translating the underlying [C string](#) to a Unicode string, making it a costly call.

In addition, you should not call getString() on primitive numeric types, like [byte](#) or int, unless it is absolutely necessary. It is much faster to call getInt() on an integer column, for example.

Avoid data type conversions

TimesTen instruction paths are so short that even small delays due to data conversion can cause a relatively large percentage increase in transaction time.

Use the appropriate getXXX() method on a ResultSet object for the data type of the data in the underlying database. For example, if the data type of the data is DOUBLE, to avoid data conversion in the JDBC driver you should call getDouble(). Similarly, use the appropriate setXXX() method on the PreparedStatement object for the input parameter in an SQL statement. For example, if you are inserting data into a CHAR column using a PreparedStatement, you should use setString().

Q 12: What should be strategy to minimise the impact on external system which is using our table's data for processing?

Answer : Best strategy should be using database views to minimize the impact on external system, If we are changing our table structure or columns and if we are able to populate same data using View then there should not be any impact on external system and they will not require any change.



View my cor

Jobs fr[Java Dev](#)[Happiest Mi](#)[Java Dev](#)[AEQUALIS](#)[Java Dev](#)[GAC Pvt Ltr](#)[Software I](#)[Solvosoft - \(](#)[Core Java](#)[Triadss Tec](#)[Nadu](#)[Java Dev](#)[Innoppl Tec](#)[Freshers v](#)[IT Compani](#)[Software I](#)[3Edge Solut](#)[Java & an](#)[sortkeys so](#)[JAVA DE'](#)[Cynosure -](#)[1 2 3 4 Ne](#)[What: title,](#)[Find Job](#)**Java Jobs**

[Opening for ".Net Developer" for one...](#)
Hamilton, NJ

[JAVA Entwickler \(SENIOR\)](#)
München, Bayern, Germany
Michael Bailey Associates - Munich

[Software Developer, C# .NET...](#)
Oxford, Oxfordshire, United Kingdom
Haybrook IT Resourcing

[Senior Java Developer](#)
Limerick, Limerick, Ireland
Morgan McKinley

[Hardware Development Security Engineer](#)
Cambridge, Cambridgeshire, United...
LA International Computer Consultants...

[POST A JOB >](#)

POWERED BY JOBTTHREAD

**DB Jobs**

[Senior Programmer/Analyst](#)
Cedarburg, WI
Systems Search

[BI/Reporting Lead](#)
Dublin, Dublin, Ireland
Morgan McKinley

[Test Engineers](#)
Bentonville, AR
SAGE IT

[Principal Full Stack Java Developer](#)
Dublin, Dublin, Ireland
Harvey Nash Group

[SAS Developer](#)
Albany, NY

[POST A JOB >](#)

POWERED BY JOBTTHREAD

**Testing job**

[SharePoint Product Team Leader](#)
Stansted, Essex, United Kingdom
Gamit

[Python Developer](#)
Dublin, Dublin, Ireland
Morgan McKinley

[C# Developer - Aylesbury Bucks](#)
Aylesbury, Buckinghamshire, United...
MAPP Ltd.

[Research & Development Test Lead](#)
Galway, Galway, Ireland
Manpower UK


[Web Search Evaluator -...](#)
Anywhere in Australia, New South...
APPEN

[POST A JOB >](#)

POWERED BY JOBTTHREAD

**Understanding Spring MVC**

Spring MVC is mostly used with Spring Web for any web application development. It is very powerful and nice layered architecture for flow and c...

**Java**

Java is a platform-independent, object-oriented programming language. It is used for developing applications that can run on any platform. Java is a very powerful and nice layered architecture for flow and c...

You might also like**facebook****Create a facebook profile** [Sign up for free today!](#)[Sign Up](#)

Posted by Java Interview Questions

[G+1](#) Recommend this on Google

Labels: database, investment bank, java investment bank, jdbc, Senior java developer interviews

[Duplicate](#)[Las vegas hotel offers](#)[Synonyms](#)[Submit](#)[Photo](#)[Followers](#)[Data Structure](#)[Search](#)[Difference](#)[2 Face](#)[Duplicate](#)[Las vegas hotel offers](#)[Synonyms](#)[Submit](#)[Photo](#)[Followers](#)[Data Structure](#)[Search](#)[infolinks](#)**No comments:****Post a Comment**

Enter your comment...


Comment as: Unknown (Goc ▾)

Sign out

☐ Notify me

Publish

Preview

[Newer Post](#)[Home](#)[Older Post](#)[Subscribe to: Post Comments \(Atom\)](#)

10 seconds test - what do you see?

● **Brain** ● **Intestine** ● **Fats**

yourtest-india

GOOGLE ANALYTICS

Picture Window template. Template images by Bim. Powered by Blogger.